

# Intro to dplyr

## Contents

<b>Data Manipulation (data cleaning)</b>	<b>1</b>
Filtering cases with <code>filter()</code> . . . . .	1
Adding new variables with <code>mutate()</code> . . . . .	2
Renaming variable with <code>rename()</code> . . . . .	2
Recoding with <code>case_when()</code> . . . . .	3
Sorting with <code>arrange()</code> . . . . .	3
Selecting variables with <code>select()</code> . . . . .	5
Descriptive statistics with <code>summarize()</code> . . . . .	6
Grouping Data with <code>group_by()</code> . . . . .	6
Grouped descriptives . . . . .	6
Aggregating Variables . . . . .	6
Pipelines . . . . .	7

---

```
acitelli <- read.csv("acitelli.csv")
```

## Data Manipulation (data cleaning)

We'll use the package `dplyr`. The `dplyr` package contains the functions for all of the data cleaning verbs: `filter()`, `mutate()`, `rename()`, `arrange()`, `select()`, `summarize()`, and `group_by()`. You can find a cheat sheet for `dplyr` [here](#).

```
#install.packages("dplyr")
library(dplyr)
```

### Filtering cases with `filter()`

First, let's filter cases. We can make a dataset of men only. Notice that we used a double equal sign, `==`, instead of single, `=`. When you want to ask if something is equal to some value or another variable, that is, you want to use equal in a *logical statement*, you need the double equal. You can also use `>` `<` `>=` `<=` `&`, which means **AND**, and finally, `|` which means **OR**.

```
menOnly <- filter(acitelli, gender == 1)
```

We could also use the pipe, `%>%`.

```
menOnly <- acitelli %>%  
  filter(gender ==1)
```

We can save this new data set in our files as a csv. This code will save in the same directory where your .Rmd file is saved. You could give a more specific file path.

```
write.csv(menOnly, "men.csv")
```

How about only the men who are above the median for `Yearsmar`. First, find the median years married.

*#use a function in the mosaic package*

Then, filter for men above that cut off point.

```
mature_hus <- menOnly %>%  
  filter(Yearsmar > -1.089)
```

Instead of first finding the median with `favstats()`, we could ask for the median inside of `filter()` with the Base R function, `median()`. Base R has all of the descriptive stats functions you'd expect, `mean()`, `sd()`, `cor()`, but be careful because if you have missing data you'll have to add `rm.na = TRUE` as an argument to the function. The syntax also differs from `mosaic`.

```
mature_hus <- menOnly %>%  
  filter(Yearsmar > median(Yearsmar))
```

## Adding new variables with `mutate()`

Let's add a new categorical variable that marks the median split on `Yearsmar`. After you create it, take a look at it.

```
menOnly <- menOnly %>%  
  mutate(mature_hus = Yearsmar > median(Yearsmar))
```

How would you get the frequencies on this variable?

*#frequencies*

Now for a sanity check, how would you get the descriptive stats split by mature and non-mature husbands?

*#descriptives split*

## Renaming variable with `rename()`

We copy a variable and give it a new name with a function you already know, `mutate()`.

```
menOnly <- menOnly %>%
  mutate(old_hus = mature_hus)
```

We can rename a variable without creating a new one with `rename(new_name = old_name)`. This is handy if you forget to name variables in Qualtrics!

```
menOnly <- menOnly %>%
  rename(wise_hus = old_hus)
```

We can rename a bunch at the same time. This is handy if you forget to name variables in Qualtrics!

```
menOnly <- menOnly %>%
  rename(self_positivity = self_pos,
         other_positivity = other_pos)
```

## Recoding with `case_when()`

Let's say we want to take gender, which is currently effects coded (men = 1 and women = -1) and make it a dummy variable. We can use the `ifelse()` function in combination with `mutate()` to achieve this.

```
acitelli <- acitelli %>%
  mutate(man = ifelse(gender == 1, 1, 0))
```

But what if we wanted to slice up years married to create a string variable that indicated newlyweds, early marriage, and mature marriages? We could nest `ifelse()` statements, but a better idea is to use the `case_when()` function.

```
acitelli <- acitelli %>%
  mutate(married = ifelse(Yearsmar < -7, "newlywed",
                        ifelse(Yearsmar > -7 & Yearsmar < 0, "early marriage", "mature marriage"))
```

These nested `ifelse()` functions can get out of control. So `case_when()` to the rescue.

```
acitelli <- acitelli %>%
  mutate(married = case_when(Yearsmar < -7 ~ "newlywed",
                            Yearsmar > -7 & Yearsmar < 0 ~ "early marriage",
                            Yearsmar >= 0 ~ "mature marriage"))
```

## Sorting with `arrange()`

First, you should know that you can sort in the viewer by clicking the (faint) arrows just to the right of each variable name. Give it a try. It's often handy to have a sort command in your code, and/or you might want to sort by more than one variable.

```
head(acitelli)
```

```
##   cuplid  Yearsmar gender self_pos other_pos satisfaction tension simhob
## 1     3   8.202667    -1     4.8     4.6     4.000000     1.5     0
## 2     3   8.202667     1     3.8     4.0     3.666667     2.5     1
## 3    10  10.452667    -1     4.6     3.8     3.166667     4.0     0
## 4    10  10.452667     1     4.2     4.0     3.666667     2.0     0
## 5    11  -8.297333    -1     5.0     4.4     3.833333     2.5     0
## 6    11  -8.297333     1     4.2     4.8     3.833333     2.5     0
##   man      married
## 1   0 mature marriage
## 2   1 mature marriage
## 3   0 mature marriage
## 4   1 mature marriage
## 5   0      newlywed
## 6   1      newlywed
```

Say we want to take a peak at the women with the bottom 6 `self_pos` scores.

```
acitelli %>%
  arrange(gender, self_pos) %>%
  head()
```

```
##   cuplid  Yearsmar gender self_pos other_pos satisfaction tension simhob
## 1    160  8.7026667    -1     3.2     3.8     3.333333     4.0     0
## 2     52 13.1193333    -1     3.4     3.8     3.833333     2.0     1
## 3    441  0.1193333    -1     3.4     4.4     4.000000     3.0     0
## 4     70 11.3693333    -1     3.6     4.4     3.833333     1.5     0
## 5    116  4.7860000    -1     3.6     4.2     2.333333     4.0     0
## 6    178 -7.0473333    -1     3.6     3.6     2.666667     3.0     0
##   man      married
## 1   0 mature marriage
## 2   0 mature marriage
## 3   0 mature marriage
## 4   0 mature marriage
## 5   0 mature marriage
## 6   0      newlywed
```

We could also save the arranged dataset.

```
acitelli <- acitelli %>%
  arrange(gender, self_pos)

head(acitelli)
```

```
##   cuplid  Yearsmar gender self_pos other_pos satisfaction tension simhob
## 1    160  8.7026667    -1     3.2     3.8     3.333333     4.0     0
```

```
## 2      52 13.1193333      -1      3.4      3.8      3.833333      2.0      1
## 3     441  0.1193333      -1      3.4      4.4      4.000000      3.0      0
## 4      70 11.3693333      -1      3.6      4.4      3.833333      1.5      0
## 5     116  4.7860000      -1      3.6      4.2      2.333333      4.0      0
## 6     178 -7.0473333      -1      3.6      3.6      2.666667      3.0      0
##  man      married
## 1      0 mature marriage
## 2      0 mature marriage
## 3      0 mature marriage
## 4      0 mature marriage
## 5      0 mature marriage
## 6      0      newlywed
```

What about the top 6? We can use the `desc()` function inside of `arrange()`.

```
acitelli %>%
  arrange(gender, desc(self_pos)) %>%
  head()
```

```
##  cuplid  Yearsmar gender self_pos other_pos satisfaction tension simhob
## 1      11 -8.297333      -1      5      4.4      3.833333      2.5      0
## 2      98 -9.214000      -1      5      4.2      4.000000      2.0      1
## 3     114 12.619333      -1      5      3.4      3.666667      2.5      0
## 4     127  3.619333      -1      5      4.6      3.833333      2.0      0
## 5     135  7.786000      -1      5      5.0      4.000000      1.5      0
## 6     177 11.619333      -1      5      5.0      4.000000      1.0      1
##  man      married
## 1      0      newlywed
## 2      0      newlywed
## 3      0 mature marriage
## 4      0 mature marriage
## 5      0 mature marriage
## 6      0 mature marriage
```

## Selecting variables with `select()`

Save a smaller subset of variables.

```
small <- acitelli %>%
  select(cuplid, gender, satisfaction, self_pos)
```

We can also save everything but some variable(s).

```
no_tension <- acitelli %>%
  select(-tension)
```

## Descriptive statistics with summarize()

```
acitelli %>%  
  summarize(mean = mean(satisfaction),  
            sd = sd(satisfaction),  
            min = min(satisfaction))
```

```
##      mean      sd      min  
## 1 3.60473 0.4964205 1.166667
```

## Grouping Data with group\_by()

### Grouped descriptives

We can split the file and view results grouped by some variable.

```
acitelli %>%  
  group_by(gender) %>%  
  summarize(mean = mean(satisfaction),  
            sd = sd(satisfaction),  
            min = min(satisfaction))
```

```
## # A tibble: 2 x 4  
##   gender mean      sd      min  
##   <int> <dbl> <dbl> <dbl>  
## 1     -1  3.59 0.530  1.50  
## 2      1  3.62 0.462  1.17
```

### Aggregating Variables

You can use `group_by()` to create aggregated variables, this is handy if you have nested data. We actually do have married couples here, so let's create a dyad mean tension variable.

```
acitelli <- acitelli %>%  
  group_by(cuplid) %>%  
  mutate(tension_mean = mean(tension)) %>%  
  ungroup()
```

*#this last command is not entirely nessesary, but good practice*

## Pipelines

We now seen our first pipelines, using `group_by()`. Now we can make a pipeline of many of the commands I did above. The last thing I do is drop useless `gender` variable, because the resulting dataset if all men.

```
mature_hus2 <- acitelli %>%  
  filter(gender == 1) %>%  
  mutate(wise_hus = Yearsmar > median(Yearsmar)) %>%  
  rename(self_positivity = self_pos,  
         other_positivity = other_pos,  
         personID = cuplid) %>%  
  arrange(wise_hus) %>%  
  select(-gender)
```

Save a dataset of women who are perceiving above the mean tension, and drop the `simhob` variable.

```
#above the mean
```

What are the couple ID's of the couples with the lowest 3 average satisfaction scores?

```
#3 lowest
```